

Rstats

Data Analysis with Rust

Presented at Brisbane Rust Group meeting on 28th July 2022

Libor Spacek

July 3, 2022

Overview

Usage

Into Cargo.toml file:

```
[dependencies]
```

```
rstats = "^1"
```

Sructs:

```
use rstats::{Mstats};
```

```
use indxvec::{MinMax};
```

```
use medians::Med;
```

Traits:

```
use rstats::{Stats, Vecg, Vecu8, MutVecg, VecVec, VecVecg};
```

Testing and Benchmarking

It is recommended to read and run `tests/tests.rs`, with examples of usage.

To run all the tests, use single thread in order to produce the results in the right order:

```
cargo test --release -- --test-threads=1 --nocapture  
--color always
```

Timing of `gmedians` uses dev dependency:

```
devtimer = "^4"
```

Introduction

Rstats is primarily about characterising multidimensional sets of points (vectors), with applications to Machine Learning and Data Analysis.

Basic statistical measures and vector algebra provide self-contained tools for the multidimensional algorithms but can also be used in their own right. Non analytical statistics is used, whereby the 'random variables' are replaced by vectors of real data. Probabilities densities and other parameters are always obtained from the data, not from some assumed distributions.

Our treatment of multidimensional sets of points is constructed from the first principles. Some original concepts, not found elsewhere, are introduced and implemented here:

New Concepts

- ▶ `median correlation`- in one dimension (1d), our `mediancorr` method is to replace *Pearson's correlation*. We define *median correlation* as cosine of an angle between two zero median vectors (instead of Pearson's zero mean vectors).
- ▶ `gmedian` - fast multidimensional geometric median (gm) algorithm.
- ▶ `madgm` - generalisation of robust data spread estimator known as 'MAD' (median of absolute deviations from median), from 1d to nd.
- ▶ `contribution` - of a point to an nd set. Defined as gm displacement when the point is added/removed. Related to the point's radius but not the same, as it depends on all the points.
- ▶ `comediance` - instead of covariance (matrix). It is obtained by supplying `covar` with the geometric median instead of the usual centroid. Thus *zero median vectors* are replacing *zero mean vectors* in covariance calculations.

Contribution

Key question of Machine Learning (ML) is how to quantify the contribution that an example point (typically a member of some large n -dimensional set) makes to the recognition concept, or outcome class, represented by that set.

In answer to this, we define `contribution` of a point as the change to `gm` caused by adding/deleting that point. Generally, more outlying points make greater contributions. However, `gm` is far less affected by outliers than centroid.

`Contribution` is computed efficiently from an existing `gm`.

Impacts of Geometric Median

Zero median vectors are generally preferable to the commonly used zero mean vectors.

In n dimensions (nd), many authors 'cheat' by using *quasi medians* (1-d medians along each axis). Quasi medians are a poor start to stable characterisation of multidimensional data. In a highly dimensional space, they are not even any faster to compute.

Specifically, all such 1d measures are sensitive to the choice of axis and thus are affected by rotation.

In contrast, analyses based on the true geometric median (gm) are axis (rotation) independent. Also, they are more stable, as medians have a 50% breakdown point (the maximum possible). They are computed here by methods `gmedian` and its weighted version `wgmedian`, in traits `vecvec` and `vecvecg` respectively.

Implementation

The main constituent parts of Rstats are its generic traits. All data are Vecs of arbitrary length d (dimensionality). Vecs give fast random access. The different traits are determined simply by how many Vecs their methods handle: 1, 2, or n .

Statisticians call them 'samples of random variables $\mathbf{X}, \mathbf{Y}, \dots$ '. Vector algebraists call them vectors $\mathbf{a}, \mathbf{b}, \dots$. Set theoreticians call them sets S, T, \dots etc. **All of these are just the same lists of numbers!** Based on this important abstraction, Rstats is able to combine numerous methods from all three branches of mathematics, plus information theory, in one relatively small and conceptually consistent crate.

The Rstats Traits

- ▶ Stats: a single vector (of numbers)
- ▶ Vecg: methods of vector algebra and information theory operating on two vectors, e.g. scalar product
- ▶ MutVecg: some of the above methods, mutating self
- ▶ Vecu8: a few methods implemented more efficiently for u8
- ▶ VecVec: methods operating on n vectors in d dimensions
- ▶ VecVecg: as VecVec but take another generic argument, such as a vector of weights.

End Types

Vectors as arguments to most methods are defined over generic type parameter `T`. Thus the `Vec` items containing the actual data can be of any primitive numeric type.

Everything will work on user defined types as well, as long as `PartialOrd`, `Copy` and conversions to/from `f64` are implemented for them.

End type `f64` is most commonly used for the results, for reasons of maintaining accuracy.

Documentation

For more detailed comments, plus some examples, see the source. You may have to unclick the 'implementations on foreign types' somewhere near the bottom of the page in the rust docs to get to it. Since these traits are implemented over the pre-existing Rust Vec type.

Reflections on Rust in Science

- ▶ + Speed and Parallelism
- ▶ + Practically no runtime errors
- ▶ - Invariance of difficulty: now lots of compiler errors
- ▶ - Unable to print any Vecs. Had to implement it myself in `indxvec::Printing`. This is not something that beginners should have to do. It is non-trivial for generic Vecs
- ▶ - Generic types: surprisingly hard battle against the grain just to use more than one standard numeric end type in Vecs
- ▶ - Turbofish needed in unpredictable ways (type inference bug?),
- ▶ - e.g. when `gm` is of concrete type `&[f64]`, instead of the expected `&[T]`, we have to annotate: `s.vsub::<f64>(gm)`.